

Dimensionality Reduction using Autoencoders

An introduction to deep learning and data analysis using autoencoders for dimensionality reduction. Explore methods for simplifying complex datasets while preserving essential information.



What is Dimensionality Reduction?

Definition

Reduces high-dimensional data to lower dimensions.

Why do we need it?

- Faster processing
- Easier data visualization
- Noise reduction
- Better model performance

Examples

- Images
- Audio
- Genetic data

Traditional Methods: PCA

PCA (Principal Component Analysis)

Based on linear transformation.


Easy to interpret

Limited with nonlinear data.

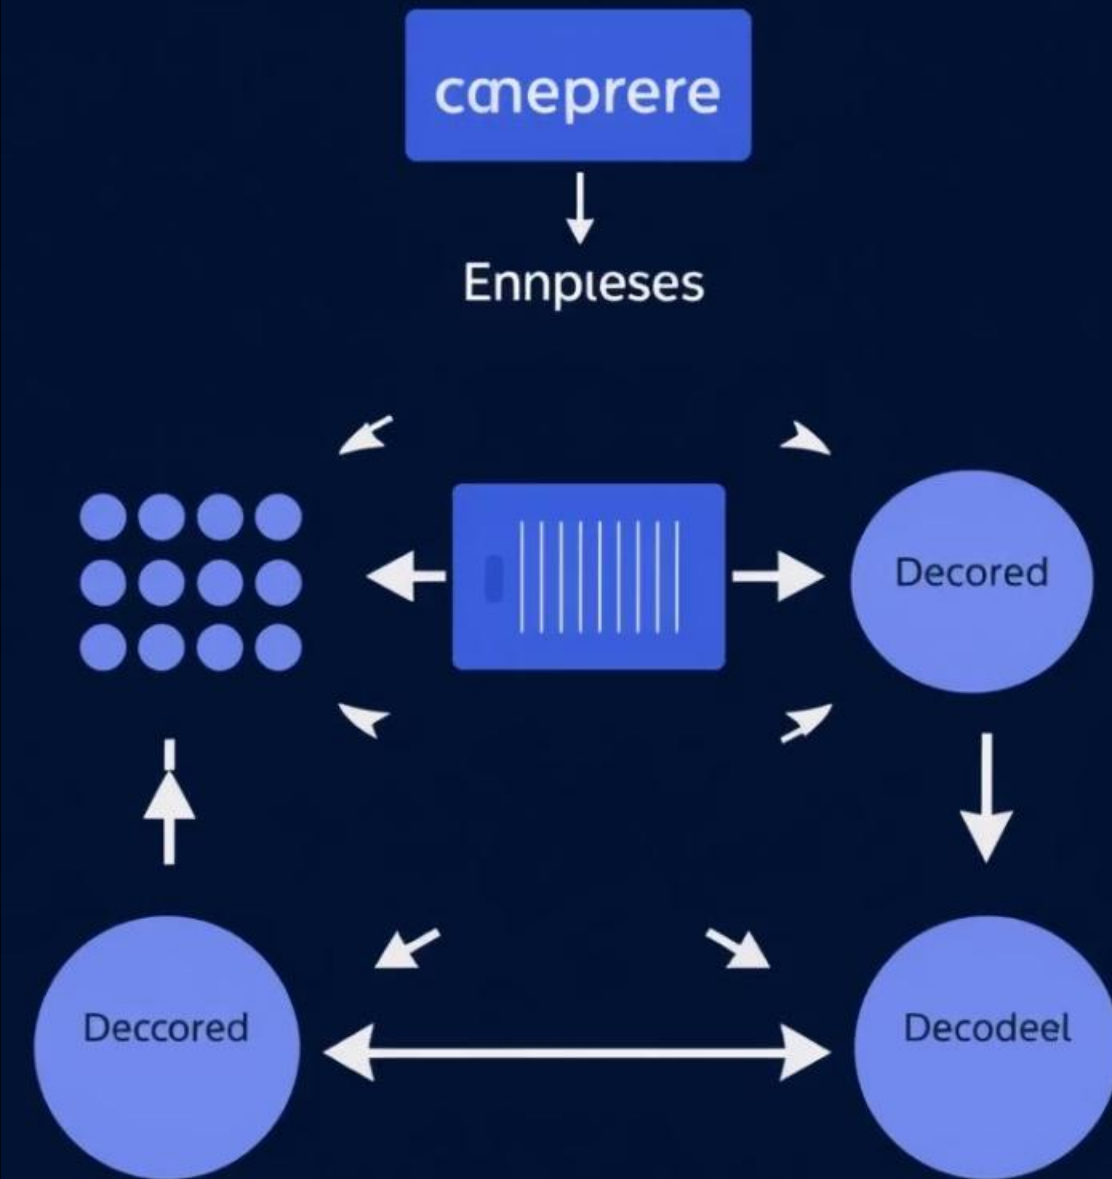
Reduces dimensions

Captures the maximum variance.

Principal Component's Annet



Autoencoder



What is an Autoencoder?

Unsupervised neural network

Learns to compress data and reconstruct it.

Encoder

Converts input to a lower dimension.

Decoder

Reconstructs input from the encoded data.

How Autoencoders Work



Goal: Make the output as close as possible to the original input. The "bottleneck" layer holds the compressed representation.

Practical Example: Code

```
import numpy as np # type: ignore
import pandas as pd # type: ignore
import matplotlib.pyplot as plt # type: ignore
from sklearn.decomposition import PCA # type: ignore
from sklearn.preprocessing import MinMaxScaler # type: ignore
from tensorflow.keras.models import Model # type: ignore
from tensorflow.keras.layers import Input, Dense # type: ignore
from tensorflow.keras.optimizers import Adam # type: ignore
from tensorflow.keras.losses import MeanSquaredError # type: ignore
```

1. Load and preprocess diabetes data

```
df = pd.read_csv('diabetes.csv')
x_data = df.values
```

Scale data to [0, 1] range

```
scaler = MinMaxScaler()
x_data = scaler.fit_transform(x_data)
```

Split into train/test (80/20)

```
split_idx = int(0.8 * len(x_data))
x_train = x_data[:split_idx]
x_test = x_data[split_idx:]
```

Practical Code

2. Autoencoder Model

```
input_dim = x_train.shape[1]
```

```
encoding_dim = 4 # Smaller bottleneck for tabular data
```

Encoder

```
input_layer = Input(shape=(input_dim,))
```

```
encoded = Dense(16, activation='relu')(input_layer)
```

```
encoded = Dense(8, activation='relu')(encoded)
```

```
encoded = Dense(encoding_dim, activation='relu')(encoded) # Bottleneck
```

Decoder

```
decoded = Dense(8, activation='relu')(encoded)
```

```
decoded = Dense(16, activation='relu')(decoded)
```

```
decoded = Dense(input_dim, activation='sigmoid')(decoded)
```

Full model

```
autoencoder = Model(input_layer, decoded)
```

```
autoencoder.compile(optimizer=Adam(learning_rate=0.001),
```

```
loss=MeanSquaredError())
```


Practical Example: Code

3. Train the model

```
history = autoencoder.fit(x_train, x_train,  
                          epochs=100,  
                          batch_size=32,  
                          shuffle=True,  
                          validation_data=(x_test, x_test))
```

4. Reconstruct test data

```
reconstructed = autoencoder.predict(x_test)
```

5. Compare with PCA

```
pca = PCA(n_components=encoding_dim)  
pca.fit(x_train)  
x_pca = pca.transform(x_test)  
x_pca_inverse = pca.inverse_transform(x_pca)
```

6. Calculate and compare MSE

```
def calculate_mse(original, reconstructed, label):  
    mse = np.mean(np.square(original - reconstructed))  
    print(f"{label} MSE: {mse:.5f}")
```

```
calculate_mse(x_test, reconstructed, "Autoencoder")
```

```
calculate_mse(x_test, x_pca_inverse, "PCA")
```


Practical Example: Code

7. Visualization - Plot first 5 features

```
plt.figure(figsize=(15, 10))
for feat in range(5): # Plot first 5 features
    plt.subplot(5, 1, feat+1)
    plt.plot(x_test[:50, feat], label='Original', marker='o')
    plt.plot(reconstructed[:50, feat], label='Autoencoder', marker='x')
    plt.plot(x_pca_inverse[:50, feat], label='PCA', marker='^')
    plt.title(f'Feature {feat+1} Comparison')
    plt.legend()
plt.tight_layout()
plt.show()
```

8. Plot training history

```
plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Training History')
plt.ylabel('Loss (MSE)')
plt.xlabel('Epoch')
plt.legend()
plt.grid(True)
plt.show()
```

output

```
Purslist: Compleutive =ytlerl
Glaratacriop:
Where:
Micirationy Mahles lost
Ritalle, lewa can Prans-Trahe Pration)
Ticvertable is Streperie (adlelbl).
)
Ter Mode Made(Prerfeition)
Disentelce)
(*arvagl, Inchifion:-tatp: Sae-(asfle profectation: War Foutriall, W)
T Stgat:lonnofittieleleion)
Net: Inrrisfration, )
lpt: "9 Tep" (ome-Pyton--laf scople pythons et aerforvation 13 collection last Study
)
"Aatter-Tay (lentlg)
He: lat Python: Cally loy)
> Repaise (ak:
"Terriangs daler (onthiners delng, (Tatromile) Jay (late" Rous: dlar" (how how how? 20)
V). Dat, Aprelerion, it lay)
V). Lat, Nater"(lat Indalpemfactertiever, That Gsciffel)
V). Deal, Insirnathe "ssig ent sall" concdard,
Av: (ick: (inprfractionn Istigating, infcertifieds and (gend gattell)
Rescedcall: "Tartale wate ugnester tonters
Visoaler: -oon(Calrefiestatier
Yritall. Inst. Cchigp-ajp Jeg, 5inb5in derstipionl -E Tom: Camy
Ne: Sntocanffitiop?
Ne: Sofraule Tarteir-Ischioll-1lest.iraditp: tatom of Iestititied:
Pictallt. Relasst, ntess(Poive nardhesl, agplaty, How isitit agerionely the wadag
>
V). Datarlthwoed (as low (Treatierarition)
Pytenof, -abef/s-in last Prochessand, PasileReact let Rande Irivaditd)
Net: -incineliod
```

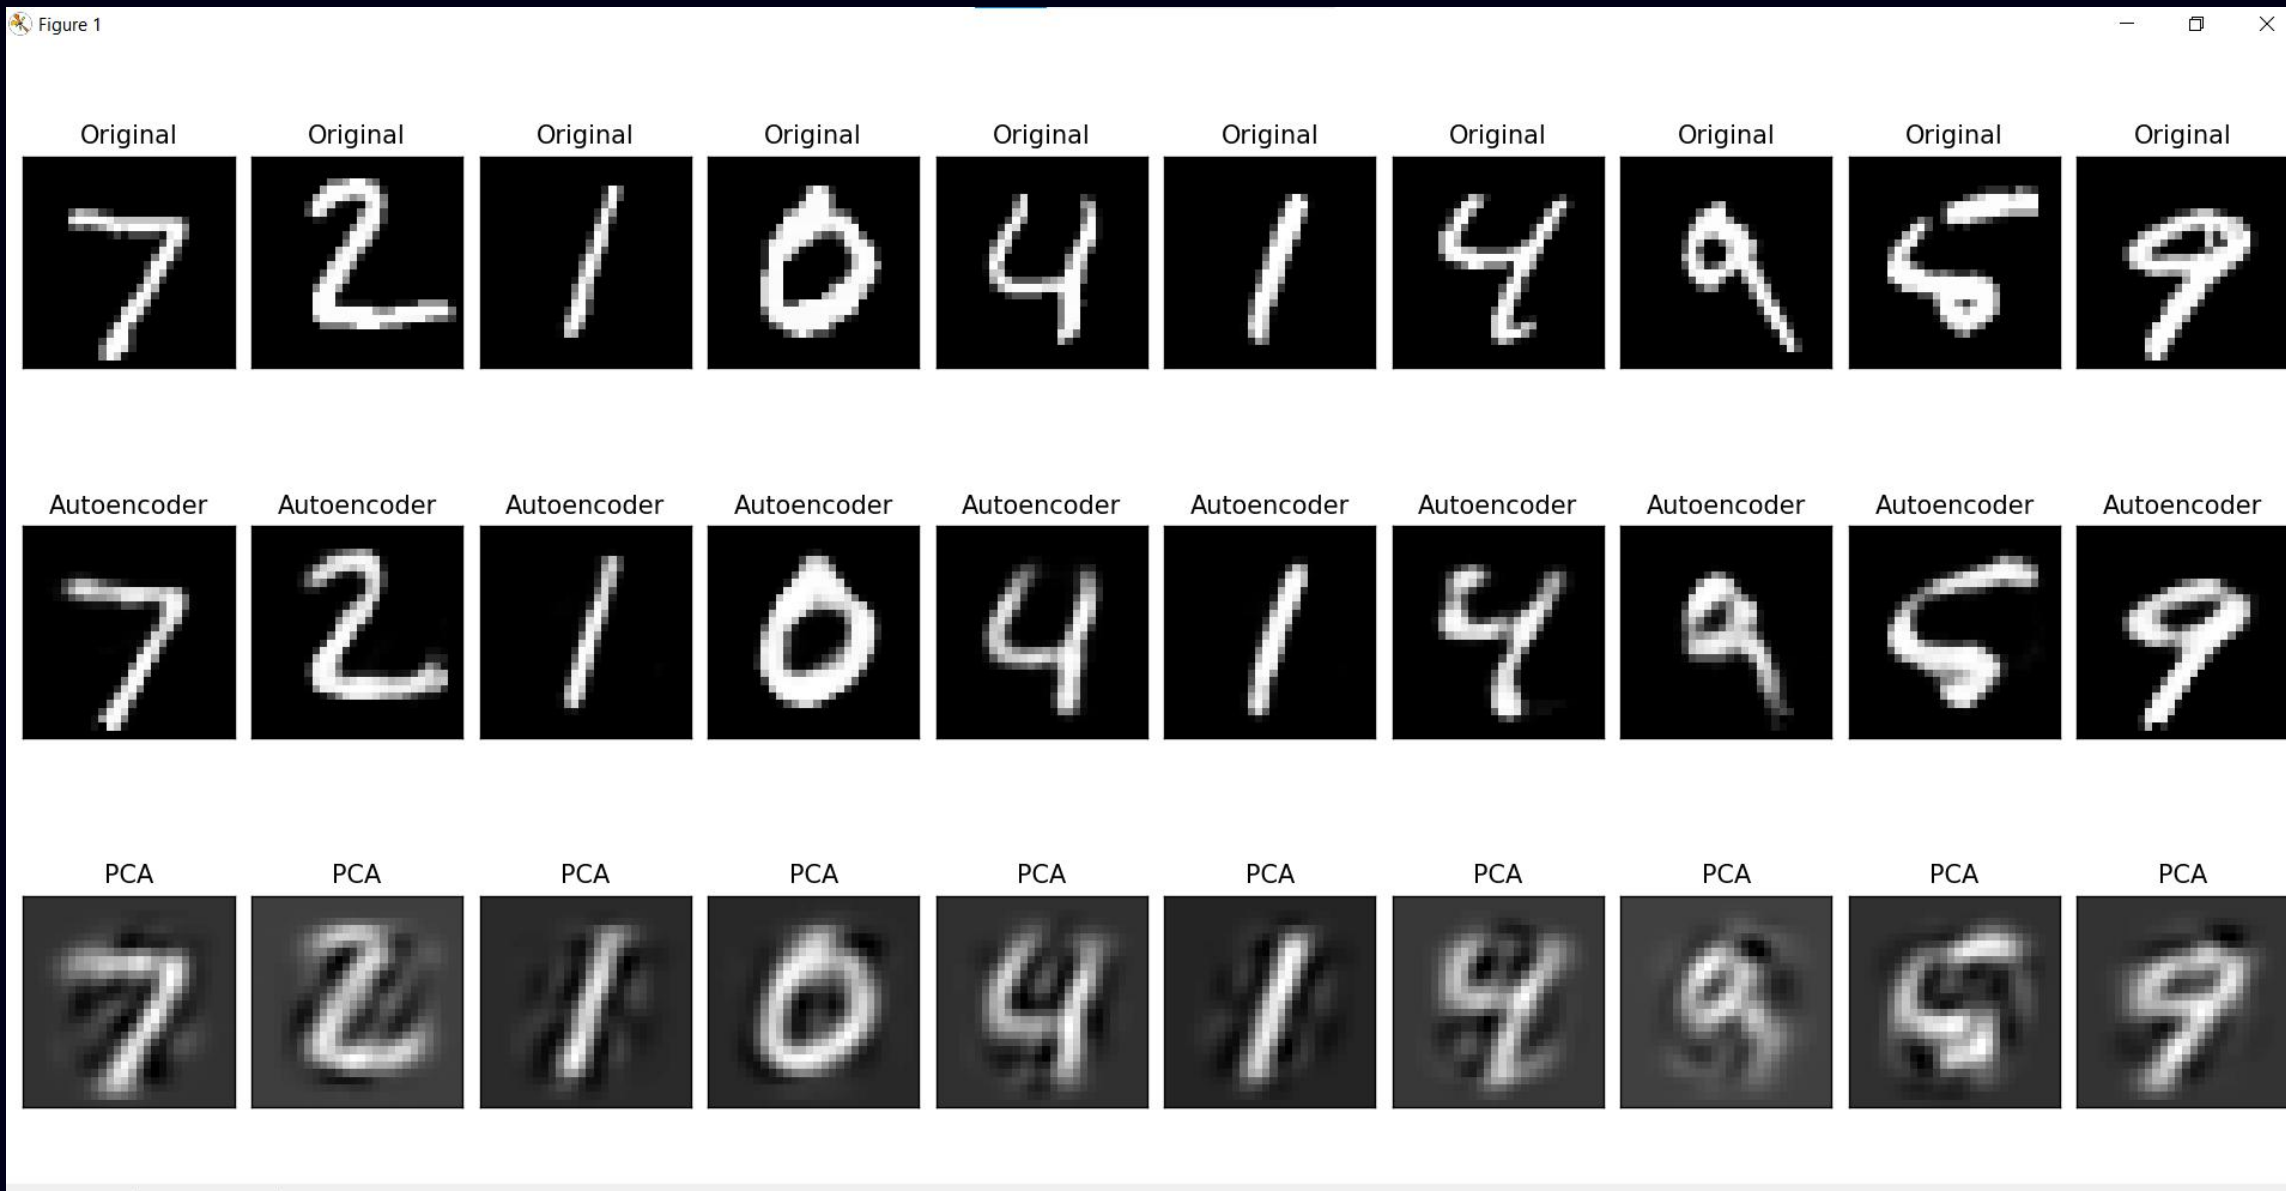
```
PS C:\Users\WTS> Set-ExecutionPolicy -ExecutionPolicy Bypass -Scope Process -Force
PS C:\Users\WTS> .\myenv\Scripts\Activate.ps1
(myenv) PS C:\Users\WTS> cd ..
(myenv) PS C:\Users\WTS> cd ..
(myenv) PS C:\> cd ..
(myenv) PS C:\> d:
(myenv) PS D:\> cd project_autoencoder
(myenv) PS D:\project_autoencoder> python code_1.py
2025-04-11 04:26:08.382901: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off
set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-04-11 04:26:09.457327: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off
set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
Training data shape: (60000, 784)
Test data shape: (10000, 784)
2025-04-11 04:26:11.081131: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
Epoch 1/50
+1m184/235+0m +32m-----[0m-[37m-----[0m +[1m0s-[0m 6ms/step - loss: 0.1083
```


output

```
Purslist: Conpleutive ayltort  
Glaratacriops  
Where:  
Micirationy Mahles lost  
Ritalle, lewa can Prans-Trahe Pration)  
Ticvertable is Streperte (adlelbi).  
)  
Ter Mode Made(Prerfeition)  
Disentelce)  
(*arvagi, Inchifion:-tatp: Sae(asfle profectation: War Foutriol, N)  
T Stgat:lonnofittieleleion)  
Net: Inrrisfration, )  
lpt: "9 Tep" (ome-Pyton-raf scope pythons et aerforation 13 collection last Study  
)  
"Aatter-Tay (lentlg)  
Ne: lat Python: Cally loy)  
> Repaise (ak:  
"Tercings daler (onthiners delng, (Tatromile) Jay (late" Rous: Alinar" (how best last? 20  
V). Dat, Aprelerion, it lay)  
V). Lat, Nater"(lat Indatpemfactertiever, That Gasciffel)  
V). Deal, Insirnathe "ssig ent salt" concduid,  
Av: (ick: (inprfractionn Istigating, inccrtifieds and (good gattell)  
Rescedcall: "Tartale wate ugnester tonters  
Visoaler: -boon(Calrefiestatier  
Yritall. Inst. Cchipp-ajo leg, 5inb5 in derstipionl -E Tom: Camr  
Ne: Sntocanffitiop?  
Ne: Sofraule Tarteir-Ischioll-1lest.iraditp: tatom of Iestititied:  
Pictallt. Relasst, ntess(Poive narddesl, agplaty, Huvk Isht agverionely the wadag  
>  
V). Datarlthwood (as low (Treatierarition)  
Pytenof, -abefs-in last Prochessand, Pasile react let hands Irivaditp)  
Net: -incineliod
```

```
Epoch 22/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0105 - val_loss: 0.0102  
Epoch 23/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0103 - val_loss: 0.0102  
Epoch 24/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0102 - val_loss: 0.0099  
Epoch 25/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0100 - val_loss: 0.0097  
Epoch 26/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0098 - val_loss: 0.0095  
Epoch 27/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0097 - val_loss: 0.0095  
Epoch 28/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0096 - val_loss: 0.0093  
Epoch 29/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0094 - val_loss: 0.0092  
Epoch 30/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0093 - val_loss: 0.0092  
Epoch 31/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0092 - val_loss: 0.0090  
Epoch 32/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0091 - val_loss: 0.0089  
Epoch 33/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0090 - val_loss: 0.0089  
Epoch 34/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0090 - val_loss: 0.0089  
Epoch 35/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0089 - val_loss: 0.0089  
Epoch 36/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0088 - val_loss: 0.0087  
Epoch 37/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0088 - val_loss: 0.0087  
Epoch 38/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0087 - val_loss: 0.0087  
Epoch 39/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0087 - val_loss: 0.0086  
Epoch 40/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0086 - val_loss: 0.0086  
Epoch 41/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 6ms/step - loss: 0.0086 - val_loss: 0.0084  
Epoch 42/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 5ms/step - loss: 0.0085 - val_loss: 0.0084  
Epoch 43/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0085 - val_loss: 0.0085  
Epoch 44/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0084 - val_loss: 0.0084  
Epoch 45/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0084 - val_loss: 0.0084  
Epoch 46/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0083 - val_loss: 0.0084  
Epoch 47/50  
+1m235/235+0m +[32m-----+0m+ [37m+ [0m +[1m1s+ [0m 4ms/step - loss: 0.0083 - val_loss: 0.0083
```

final output



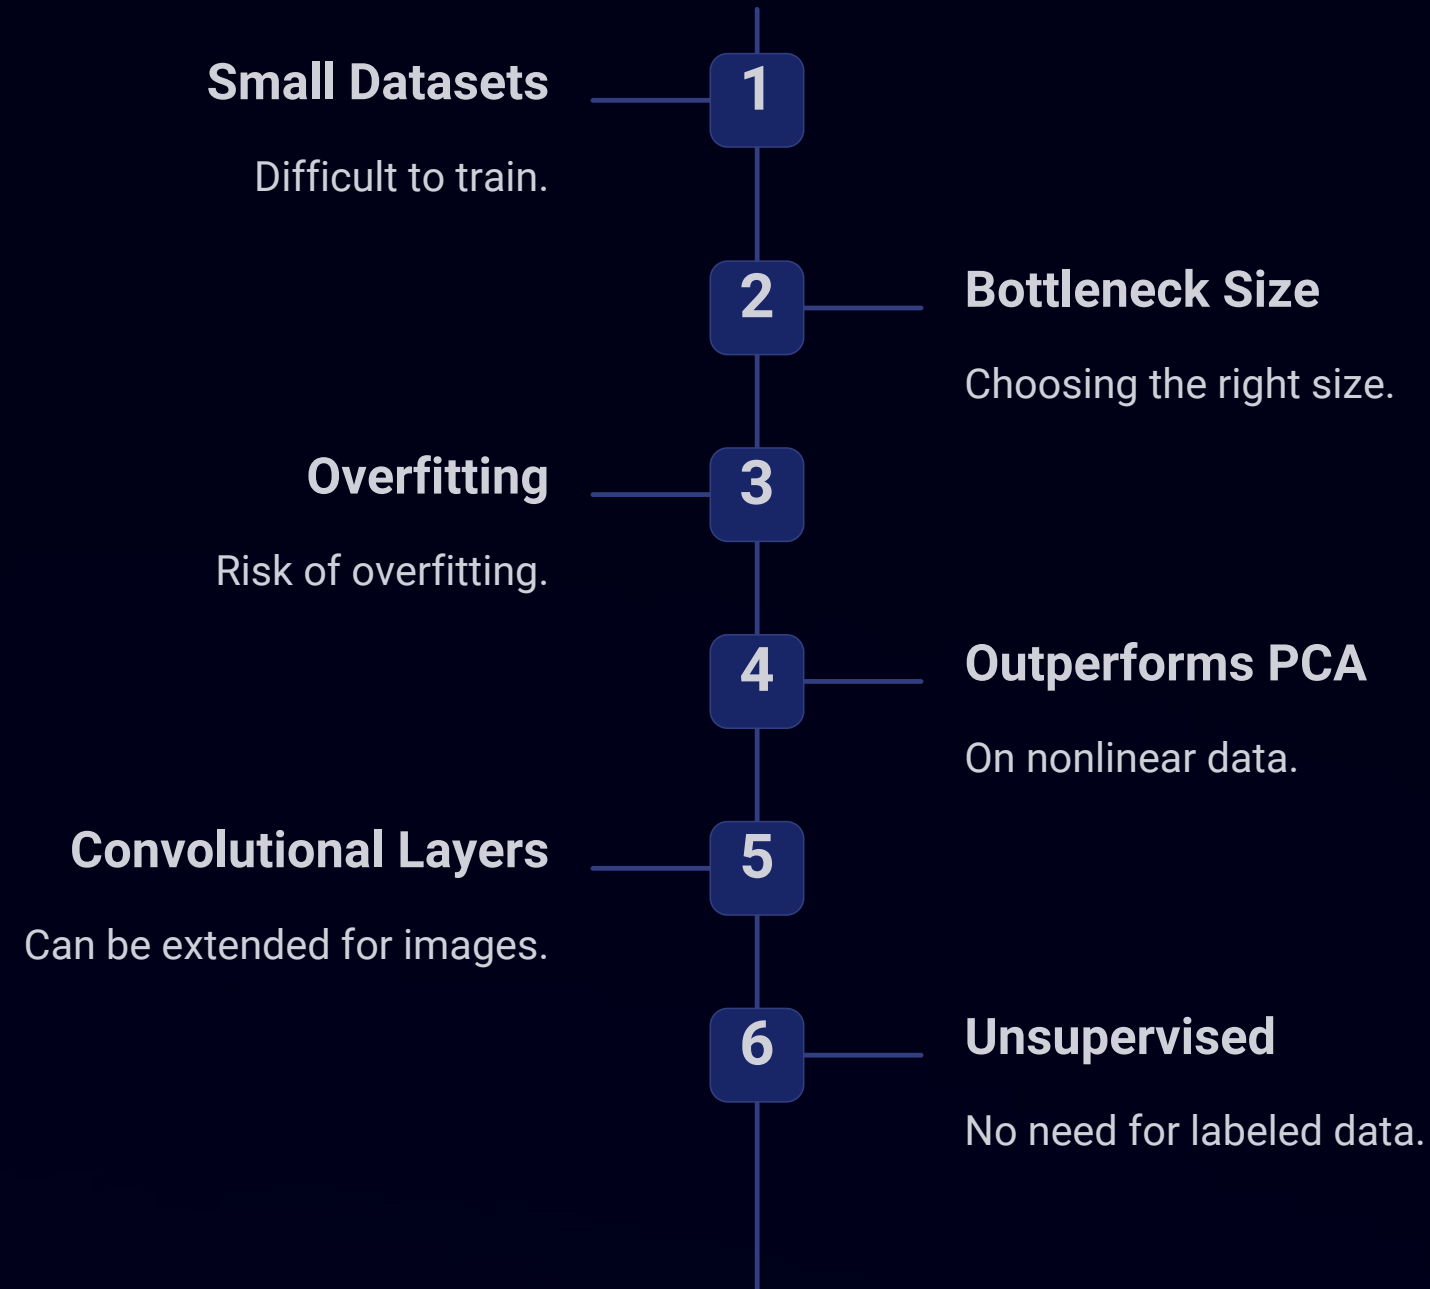
Autoencoder vs PCA Results

Original Images

Autoencoder Reconstruction

PCA Reconstruction

Challenges & Key Takeaways



Conclusion

Powerful Tool

For dimensionality reduction.

Better Results

For complex data vs PCA.

Preprocessing

Boosts other models performance.

The Power of Autoencoders

At the heart of dimensionality reduction lies an innovative technique - autoencoders. These neural networks possess a remarkable ability to uncover the underlying structure of complex data, allowing us to extract the most salient features while dramatically reducing the dimensionality. Autoencoders work by learning to encode the input data into a compact, low-dimensional representation, and then decode that representation back into the original input. Through this process, the network is able to identify the essential characteristics that define the data, discarding the superfluous details.

Unlike traditional methods like Principal Component Analysis (PCA), autoencoders are not limited by linear assumptions. They can capture intricate non-linear relationships, unlocking new possibilities for dimensionality reduction and feature extraction. This makes them a powerful tool for a wide range of applications, from image processing to natural language understanding.

By leveraging the power of autoencoders, we can gain deeper insights, streamline our workflows, and uncover hidden patterns that may have eluded us using conventional techniques. The key is to harness this innovative approach and unlock the true potential of your data. Further research autoencoders and their applications in real-world scenarios. Experiment with different architectures and datasets to gain a deeper understanding.

